

JMLR T_EX.jl Example

Robert Moss

Stanford University, Stanford, CA 94305

MOSSR@CS.STANFORD.EDU

August 21, 2020

1. Loss Function

In mathematical optimization, statistics, decision theory and machine learning, a *loss function* or *cost function* is a function that maps an event or values of one or more variables onto a real number intuitively representing some “cost” associated with the event.¹ An optimization problem seeks to minimize a loss function. An objective function is either a loss function or its negative (sometimes called a *reward function* or a *utility function*), in which case it is to be maximized.

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \left[-y_i \log(h_{\theta}(x_i)) - (1 - y_i) \log(1 - h_{\theta}(x_i)) \right] \quad (1)$$

```
1 function loss_function(theta, X, y)
2     m = length(y) # number of training examples
3     grad = zeros(size(theta))
4     h = sigmoid(X * theta)
5     J = 1/m*sum(-y'*log(h)-(1.-y)*log(1.-h))
6     grad = 1/m*(X'*(h-y))
7     return (J, grad)
8 end
```

1. https://en.wikipedia.org/wiki/Loss_function